

# Lab 4 - Part 1

Yaqi Shi, 1003813180

2023-02-04

The lab exercise results are in the end of this file and it contains only Q1 to Q3

## Extracting data on opioid prescriptions from CDC

We're going to grab some data on opioid prescription rates from the CDC website. While the data are nicely presented and mapped, there's no nice way of downloading the data for each year as a csv or similar form. So let's use `rvest` to extract the data. We'll also load in `janitor` to clean up column names etc later on.

### Getting the data for 2008

Have a look at the website at the url below. It shows a map of state prescription rates in 2008. Let's read in the html of this page.

```
cdcpage <- "https://www.cdc.gov/drugoverdose/rxrate-maps/state2008.html"
cdc <- read_html(cdcpage)
cdc
```

```
## {html_document}
## <html lang="en-us" class="cdc-2022 theme-purple cdc-page-type-content">
## [1] <head>\n<meta http-equiv="Content-Type" content="text/html; charset=UTF-8 ...
## [2] <body class="no-js cdc-page">\r\n\t<div id="skipmenu">\r\n\t\t<a class="s ...
```

Note that it has two main parts, a head and body. For the majority of use cases, you will probably be interested in the body. You can select a node using `html_node()` and then see its child nodes using `html_children()`.

```
body_nodes <- cdc %>%
  html_node("body") %>%
  html_children()
head(body_nodes)
```

```
## {xml_nodeset (6)}
## [1] <div id="skipmenu">\r\n\t\t<a class="skippy sr-only-focusable" href="#con ...
## [2] <div class="header-language-bar container text-right pt-1 pb-1 fs0875">\r ...
## [3] <header id="page_banner" role="banner" aria-label="Banner"><div class="co ...
## [4] <div class="container d-flex flex-wrap body-wrapper bg-white">\r\n\t\t<t< ...
## [5] <footer class="" role="contentinfo" aria-label="Footer"><div class="conta ...
## [6] <script src="https://www.cdc.gov/config/cdc_config.js"></script>
```

### Inspecting elements of a website

The above is still fairly impenetrable. But we can get hints from the website itself. Using Chrome (or Firefox) you can highlight a part of the website of interest (say, 'Alabama'), right click and choose 'Inspect'. That gives you info on the underlying html of the webpage on the right hand side. Alternatively, and probably easier to find what we want, right click on the webpage and choose View Page Source. This opens a new

window with all the html. Do a search for the word 'Alabama'. Now we can see the code for the table. We can see that the data we want are all within `tr`. So let's extract those nodes:

```
head(cdc %>% html_nodes("tr"))

## {xml_node_set (6)}
## [1] <tr>\n<th>State</th>\n<th>State Abbreviation</th>\n<th>Opioid Dispensing ...
## [2] <tr>\n<td>Alabama</td>\n<td>AL</td>\n<td>126.1</td>\n</tr>\n
## [3] <tr>\n<td>Alaska</td>\n<td>AK</td>\n<td>68.5</td>\n</tr>\n
## [4] <tr>\n<td>Arizona</td>\n<td>AZ</td>\n<td>80.9</td>\n</tr>\n
## [5] <tr>\n<td>Arkansas</td>\n<td>AR</td>\n<td>112.1</td>\n</tr>\n
## [6] <tr>\n<td>California</td>\n<td>CA</td>\n<td>55.1</td>\n</tr>\n
```

Great, now we're getting somewhere. We only want the text, not the html rubbish, so let's extract that:

```
table_text <- cdc %>%
  html_nodes("tr") %>%
  html_text()

head(table_text)

## [1] "State\nState Abbreviation\nOpioid Dispensing Rate per 100\n"
## [2] "Alabama\nAL\n126.1\n"
## [3] "Alaska\nAK\n68.5\n"
## [4] "Arizona\nAZ\n80.9\n"
## [5] "Arkansas\nAR\n112.1\n"
## [6] "California\nCA\n55.1\n"
```

This is almost useful! Turning it into a tibble and using `separate` to get the variables into separate columns gets us almost there:

```
rough_table <- table_text %>%
  as_tibble() %>%
  separate(value, into = c("state", "abbrev", "rate"), sep = "\n", extra = "drop")
head(rough_table)
```

```
## # A tibble: 6 x 3
##   state      abbrev      rate
##   <chr>      <chr>      <chr>
## 1 State      State Abbreviation Opioid Dispensing Rate per 100
## 2 Alabama    AL              126.1
## 3 Alaska     AK              68.5
## 4 Arizona    AZ              80.9
## 5 Arkansas   AR              112.1
## 6 California CA              55.1
```

Now we can just divert to our standard tidyverse cleaning skills (`janitor` functions help here) to tidy it up:

```
d_prescriptions <- rough_table %>%
  janitor::row_to_names(1) %>%
  janitor::clean_names() %>%
  rename(prescribing_rate = opioid_dispensing_rate_per_100) %>%
  mutate(prescribing_rate = as.numeric(prescribing_rate))

head(d_prescriptions)

## # A tibble: 6 x 3
##   state      state_abbreviation prescribing_rate
```

##	<chr>	<chr>	<dbl>
## 1	Alabama	AL	126.
## 2	Alaska	AK	68.5
## 3	Arizona	AZ	80.9
## 4	Arkansas	AR	112.
## 5	California	CA	55.1
## 6	Colorado	CO	67.7

Now we have clean data for 2008!

## Take-aways

This example showed you how to extract a particular table from a particular website. The take-away is to inspect the page html, find where what you want is hiding, and then use the tools in `rvest` (`html_nodes()` and `html_text()` particularly useful) to extract it.

## Question 1

Add a year column to `d_prescriptions`.

**Answer**

Add a year column to the dataset and fill in with value “2008”

```
d_prescriptions$year <- rep(2008, nrow(d_prescriptions))
```

## Getting all the other years

Now I want you to get data for 2008-2019 and save it into one big tibble. If you go to [cdc.gov/drugoverdose/rxrate-maps/index.html](https://www.cdc.gov/drugoverdose/rxrate-maps/index.html), on the right hand side there's hyperlinks to all the years under “U.S. State Opioid Dispensing Rate Maps”.

Click on 2009. Look at the url. Confirm that it's exactly the same format as the url for 2008, except the year has changed. This is useful, because we can just loop through in an automated way, changing the year as we go.

## Question 2

Make a vector of the urls for each year, storing them as strings.

**Answer**

First, I checked that for 2009, the url is “<https://www.cdc.gov/drugoverdose/rxrate-maps/state2009.html>” which is in the same format as the url for 2008 except that the year changes. Now we create a vector of url for each year and store it as string.

```
# Create the url lists
year_list <- 2008:2019
url_list <- rep("https://www.cdc.gov/drugoverdose/rxrate-maps/state2008.html", 12)

for(i in 1:12) {
  yr <- year_list[i]
```

```

url_list[i] <- paste("https://www.cdc.gov/drugoverdose/rxrate-maps/state", yr, ".html", sep = "")
}

# Display the results
url_list <- tibble(url = url_list, year = year_list)
url_list

## # A tibble: 12 x 2
##   url                                     year
##   <chr>                                <int>
## 1 https://www.cdc.gov/drugoverdose/rxrate-maps/state2008.html 2008
## 2 https://www.cdc.gov/drugoverdose/rxrate-maps/state2009.html 2009
## 3 https://www.cdc.gov/drugoverdose/rxrate-maps/state2010.html 2010
## 4 https://www.cdc.gov/drugoverdose/rxrate-maps/state2011.html 2011
## 5 https://www.cdc.gov/drugoverdose/rxrate-maps/state2012.html 2012
## 6 https://www.cdc.gov/drugoverdose/rxrate-maps/state2013.html 2013
## 7 https://www.cdc.gov/drugoverdose/rxrate-maps/state2014.html 2014
## 8 https://www.cdc.gov/drugoverdose/rxrate-maps/state2015.html 2015
## 9 https://www.cdc.gov/drugoverdose/rxrate-maps/state2016.html 2016
## 10 https://www.cdc.gov/drugoverdose/rxrate-maps/state2017.html 2017
## 11 https://www.cdc.gov/drugoverdose/rxrate-maps/state2018.html 2018
## 12 https://www.cdc.gov/drugoverdose/rxrate-maps/state2019.html 2019

```

### Question 3

Extract the prescriptions data for the years 2008-2019, and store in the one tibble. Make sure you have a column for state, state abbreviation, prescription rate and year. Note if you are looping over years/urls (which is probably the easiest thing to do), it's good practice to include a `Sys.sleep(1)` at the end of your loop, so R waits for a second before trying again.

Plot prescriptions by state over time.

#### Answer

First, we extract the data from the website based on what we did for 2008 data:

```

# Set up the tibble
all_data <- c()

for(i in 1:12) {
  # Extract the information
  cdcpage <- url_list$url[i]
  cdc <- read_html(cdcpage)

  # Format the data
  table_text <- cdc %>%
    html_nodes("tr") %>%
    html_text()

  rough_table <- table_text %>%
    as_tibble() %>%
    separate(value, into = c("state", "abbrev", "rate"), sep = "\n", extra = "drop")
}

```

```

d_prescriptions <- rough_table %>%
  janitor::row_to_names(1) %>%
  janitor::clean_names() %>%
  rename(prescribing_rate = opioid_dispensing_rate_per_100) %>%
  mutate(prescribing_rate = as.numeric(prescribing_rate))

colnames(d_prescriptions) <- c("state", "state_abbreviation", "prescribing_rate")

d_prescriptions <- d_prescriptions %>%
  filter(state_abbreviation != "US")

d_prescriptions$year <- rep(url_list$year[i], nrow(d_prescriptions))

# Store the results
all_data <- rbind(all_data, d_prescriptions)
}

```

Now we have extracted the data, we plot the prescriptions results for each state over time

```

ggplot(all_data, aes(x = year, y=prescribing_rate)) +
  geom_line(aes(colour=state_abbreviation)) +
  labs(title = "Prescriptions results for each state over time", x = "Year", y = "Prescribing Rate")

```

